

CS111 F21 Homework 3

Michael Glushchenko, 9403890

October 16, 2021

1 Solving $A^T x = b$

When solving $Ax = b$, we used $(PA = LU) \Rightarrow (A = P^T LU)$ to substitute:

$$Ax = b \Rightarrow (P^T LU)x = b \Rightarrow LUx = Pb$$

Then we set $y = Ux$ and solved $Ly = Pb$ by using

$$y = \text{Lsolve}(L, b[p], \text{unit_diag} = \text{True})$$

After, we got our solution by using

$$x = \text{Usolve}(U, y)$$

Now we're solving $A^T x = b$. Transpose both sides of the $LU = PA$ equation:

$$(LU = PA) \Rightarrow ((LU)^T = (PA)^T)$$

Simplify

$$\Rightarrow U^T L^T = A^T P^T \Rightarrow A^T = U^T L^T P$$

Substitute

$$A^T x = b \Rightarrow (U^T L^T P)x = b$$

Following previous logic, we set $y = L^T Px$ and use

$$y = \text{Lsolve}(U^T, b)$$

Then we set $d = Px$ and solve $L^T d = y$ via

$$d = \text{Usolve}(L^T, y, \text{unit_diag} = \text{True})$$

Vector d should simply be a permutation of the sought-after solution x :

$$d = Px \Rightarrow x = P^T d$$

In Python, we generate our random matrix A and random vector b . After calling `cs111.LUfactor(A)` we now have L, U , and p . Following the logic described above, we check that our answer is correct. Here's the code and output:

```
A = np.random.rand(6,6)
b = np.random.rand(6)
print("A: \n", A)
print("\nb: \n", b)

L, U, p = cs111.LUfactor(A)
Ptrans = np.eye(len(p))[p,:].T # create matrix P from p, tranpose it
pinv = np.array(np.where(Ptrans[:,:]==1))[1] # find pinv

y = cs111.Lsolve(U.T,b)
d = cs111.Usolve(L.T,y,unit_diag=True)
x = d[pinv] #permute the final answer
res = npla.norm(b - A.T@x) / npla.norm(b) #relative res norm
print("\nx: \n", x)
print("\nrelative residual norm =\n", res)
```

A:

```
[[0.4674 0.9857 0.3668 0.6576 0.767 0.7292]
 [0.9739 0.4032 0.0443 0.6854 0.7587 0.9184]
 [0.5428 0.4715 0.091 0.4718 0.7177 0.4702]
 [0.7252 0.8702 0.1897 0.0736 0.5873 0.8138]
 [0.84 0.3914 0.7607 0.488 0.3282 0.7324]
 [0.8315 0.2754 0.9995 0.6815 0.8485 0.4152]]
```

b:

```
[0.7463 0.592 0.3302 0.7511 0.8578 0.9784]
```

x:

```
[1.4841 2.3734 -4.057 0.4919 -1.8213 1.343]
```

relative residual norm =

```
5.384326597898874e-16
```

2 Condition Number of A

2.1

For this part, we have: $\begin{pmatrix} \alpha & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} x_0 \\ x_1 \end{pmatrix} = \begin{pmatrix} 1 - \alpha \\ 0 \end{pmatrix}$ for some $\alpha < 1$.

The code and output for this part are as follows:

```
alpha = np.array([10.0**(-4.0), 10.0**(-8.0),
                  10.0**(-12.0), 10.0**(-16.0), 10.0**(-20.0)])
x_exact = np.array([-1,1])

print("PART ONE")
for j in np.array([True, False]):
    print("\n\nPivoting = ", j, "\n\n")
    for i in alpha:
        print("alpha = ", i)
        A = np.array([[i,1],[1,1]])
        b = np.array([1-i,0])
        ans, res = cs111.LUsolve(A,b,pivoting=j)
        error = x_exact - ans
        print("condition number of A = ", npla.cond(A,2))
        print("computed x = ", ans)
        print("norm of the error = ", npla.norm(error))
        print("relative residual norm = ", res, "\n")
```

Pivoting = True	Pivoting = False
alpha = 0.0001 condition number of A = 2.618385273654827 computed x = [-1. 1.] norm of the error = 0.0 relative residual norm = 0.0	alpha = 0.0001 condition number of A = 2.618385273654827 computed x = [-1. 1.] norm of the error = 1.1013412404281553e-13 relative residual norm = 1.1014513855667119e-13
alpha = 1e-08 condition number of A = 2.618034023874506 computed x = [-1. 1.] norm of the error = 0.0 relative residual norm = 0.0	alpha = 1e-08 condition number of A = 2.618034023874506 computed x = [-1. 1.] norm of the error = 5.024759275329416e-09 relative residual norm = 5.024759325577009e-09
alpha = 1e-12 condition number of A = 2.618033988753407 computed x = [-1. 1.] norm of the error = 0.0 relative residual norm = 0.0	alpha = 1e-12 condition number of A = 2.618033988753407 computed x = [-1. 1.] norm of the error = 2.212172012150404e-05 relative residual norm = 2.212172012152616e-05
alpha = 1e-16 condition number of A = 2.6180339887498953 computed x = [-1. 1.] norm of the error = 0.0 relative residual norm = 0.0	alpha = 1e-16 condition number of A = 2.6180339887498953 computed x = [1.1102 1.] norm of the error = 2.1102230246251565 relative residual norm = 2.110223024625157
alpha = 1e-20 condition number of A = 2.6180339887498953 computed x = [-1. 1.] norm of the error = 0.0 relative residual norm = 0.0	alpha = 1e-20 condition number of A = 2.6180339887498953 computed x = [0. 1.] norm of the error = 1.0 relative residual norm = 1.0

Clearly, the matrix stays well-conditioned regardless of α or truth value of *pivoting*. When *pivoting* = *True*, the relative residual norm is 0 for all α that we tested, as our answer seems to be exact. When *pivoting* = *False*, however, the relative residual norm grows as α decreases (without *pivoting* the pivot used is very small, thus being too close to 0, thus making us close to division by 0 during *LU-factorization*, thus furthering us from the exact answer).

2.2

For this part, we have: $\begin{pmatrix} 1+\alpha & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} x_0 \\ x_1 \end{pmatrix} = \begin{pmatrix} -\alpha \\ 0 \end{pmatrix}$ for some $\alpha < 1$.

The code and output for this part are as follows:

```
alpha = np.array([10.0**(-4.0), 10.0**(-8.0),
                  10.0**(-12.0), 10.0**(-16.0), 10.0**(-20.0)])
x_exact = np.array([-1,1])

print("PART TWO")
for j in np.array([True, False]):
    print("\n\nPivoting = ", j, "\n\n")
    for i in alpha:
        print("alpha = ", i)
        A = np.array([[1+i,1],[1,1]])
        b = np.array([-i,0])
        print("condition number of A = ", npla.cond(A,2))
        try:
            ans, res = cs111.LUsolve(A,b,pivoting=j)
            error = x_exact - ans
            print("computed x = ", ans)
            print("norm of the error = ", npla.norm(error))
            print("relative residual norm = ", res, "\n")
        except:
            print("Answer not returned \n")

Pivoting = True

alpha = 0.0001
condition number of A = 40002.00007504555
computed x = [-1. 1.]
norm of the error = 2.3488583181909273e-13
relative residual norm = 1.1018204577883939e-13

alpha = 1e-08
condition number of A = 399999992.40613085
computed x = [-1. 1.]
norm of the error = 1.0153631235080597e-08
relative residual norm = 6.0774709918447105e-09

alpha = 1e-12
condition number of A = 3997936426769.096
computed x = [-0.9999 0.9999]
norm of the error = 0.00012571323468403843
relative residual norm = 2.212172012148393e-05

alpha = 1e-16
condition number of A = 5.961777047638983e + 16
Answer not returned

alpha = 1e-20
condition number of A = 5.961777047638983e + 16
Answer not returned

alpha = 0.0001
condition number of A = 40002.00007504555
computed x = [-1. 1.]
norm of the error = 2.3488583181909273e-13
relative residual norm = 1.1018204577883939e-13

alpha = 1e-08
condition number of A = 399999992.40613085
computed x = [-1. 1.]
norm of the error = 1.0153631235080597e-08
relative residual norm = 6.0774709918447105e-09

alpha = 1e-12
condition number of A = 3997936426769.096
computed x = [-0.9999 0.9999]
norm of the error = 0.00012571323468403843
relative residual norm = 2.212172012148393e-05

alpha = 1e-16
condition number of A = 5.961777047638983e + 16
Answer not returned

alpha = 1e-20
condition number of A = 5.961777047638983e + 16
Answer not returned
```

In the 3 tries where solutions were returned, we can see the relative residual norm increasing (same reason as 2.1). As α decreases, matrix A becomes *ill-conditioned* under both *pivoting* and *non-pivoting* solutions; the matrix gets closer to having $row_1 = row_2$, making one of the *pivots* very small. In trials 4 and 5, we finally get too close to dividing by 0 during *LU-factorization*. Running `cs111.LUfactor(A)` results in a non-zero pivot not being found, rendering the factorization impossible, and making `cs111.LUsolve(A,b)` crash.

3 3D Temperature Matrix

Here's my code for *make_A_3D(k)*:

```
def make_A_3D(k):
    triples = []
    for x in range(k):
        for y in range(k):
            for z in range(k):
                row = x + k*y + k*k*z # row for grid point (x,y,z)
                col = row
                triples.append((row, col, 6.0)) # diagonal elements

                # x dimension
                if x > 0:
                    col = row - 1
                    triples.append((row, col, -1.0))
                if x < k - 1:
                    col = row + 1
                    triples.append((row, col, -1.0))

                # y dimension
                if y > 0:
                    col = row - k
                    triples.append((row, col, -1.0))
                if y < k - 1:
                    col = row + k
                    triples.append((row, col, -1.0))

                # z dimension
                if z > 0:
                    col = row - k*k
                    triples.append((row, col, -1.0))
                if z < k - 1:
                    col = row + k*k
                    triples.append((row, col, -1.0))

    ndim = k*k*k # 3 dimensions
    rownum = [t[0] for t in triples]
    colnum = [t[1] for t in triples]
    values = [t[2] for t in triples]


    A = scipy.sparse.csr_matrix((values, (rownum, colnum)), shape = (ndim, ndim))
    return A
```

Starting on the next page, we go through each k in the range $[2, 3, 4, 5]$ by using the following code:

```
for k in range(2,6):
    A = make_A_3D(k)
    print("k:", k)
    print("dimensions:", A.shape)
    print("nonzeros:", A.nnz)
    #print("A as sparse matrix:\n", A)
    print("A as dense matrix:\n", A.toarray())
    plt.spy(A)
```

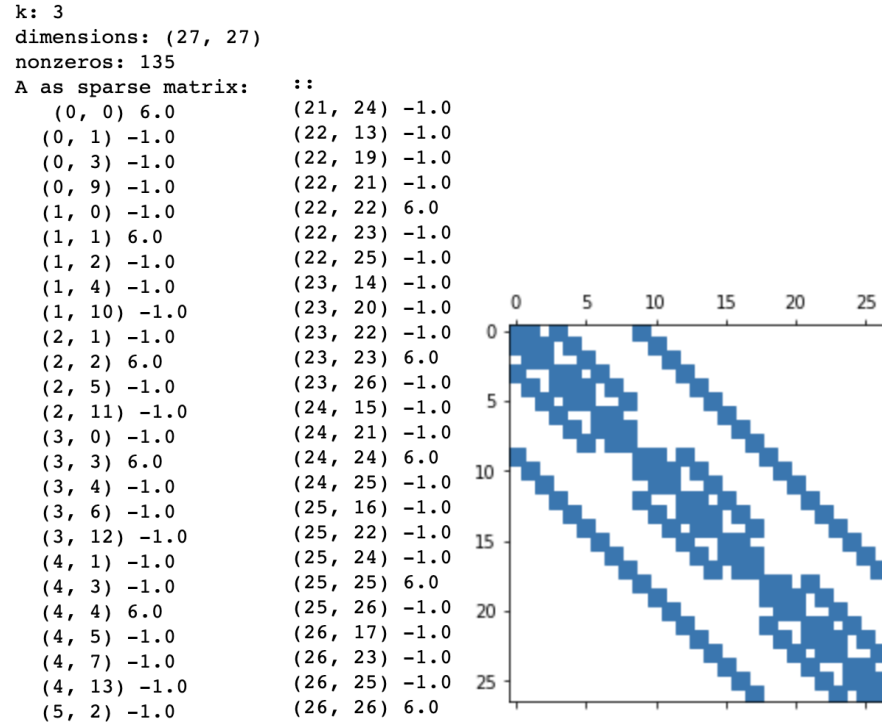
3.1 $k = 2$

Here, we get the following dense matrix A and chart:



3.2 $k = 3$


Here, we look at the sparse matrix A , since the dense matrix becomes too large:



3.3 $k = 4$

```
x_min = 0
x_max = 30
y_min = 30
y_max = 0
plt.axis([x_min, x_max, y_min, y_max])
```

Zooming in produces the following plots (zoomed chart on the right):



3.4 $k = 5$

Now, we adjust the coordinates and inspect $k = 5$:

